

BERTPlay - A semantic similarity checking project

- By: *Ramakrishnan Sundareswaran* (Student ID : 897255379, email : ramkris@iastate.edu),
Kumara Sri Harsha Vajjhala (Student ID : 571966414 , email : harshavk@iastate.edu)

Abstract

In this project, we build and experiment with different deep learning models to predict the semantic similarity between two sentences. From our experiments, we have observed that the transformer based BERT models outperformed the LSTM network and also trained faster. Moreover, the BERT models that were finetuned with a Bi-directional LSTM gave a much better performance than the BERT model that was not finetuned to the specific task of semantic similarity recognition. We have developed two games, SemanticSimilarityGame and ReadingComprehensionGame, both of which are fun to play and a good way to learn English language.

Introduction

Semantic similarity checking is an important application, with usecases in the areas of language understanding, question answering and plagiarism detection. The reason we were drawn towards working on a semantic similarity based game was because it involved creating a novel way to learn language, and also provided us with an opportunity to tinker with various state of the art NLP models

Methodology

Semantic similarity prediction task requires a deep understanding of the language, and hence, we require complex models which create word embeddings based on the context in which it is used. We explore two of the most successful architectures in NLP: BERT, and a LSTM and compare their performance.

The common dataset used to train all the models is : Stanford Natural Language Inference (SNLI). We have also experimented with a BERT-large pre-trained on SQuAD (Stanford Question Answering Dataset)

First, we clean the SNLI dataset and divide it into train, validation and test data that can be used in the different models that we are testing. For each of the BERT models, we first create a BERT tokenizer that basically encodes the given sentence with the pre-trained BERT embeddings and

adds the CLS, SEP embeddings to distinguish between the Sentence1 and Sentence2. Then, we load and train just the output layer (Dense) with the train data. In case we are finetuning, we train the model end-to-end, including the BERT. The model obtained from the above steps is then tested.

LSTM model is trained in a similar way, by encoding using word2vec, and then training the LSTM model built for 20 epochs.

We have used tensorflow for the output layer models, huggingface transformers library for the BERT models, and numpy and pandas for data cleaning.

We have used the colab notebook by Mohammed Merchant :

https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/nlp/ipynb/semantic_similarity_with_bert.ipynb#scrollTo=w60qwFa1dWE as a reference while creating our BERT models.

We have used the github code :

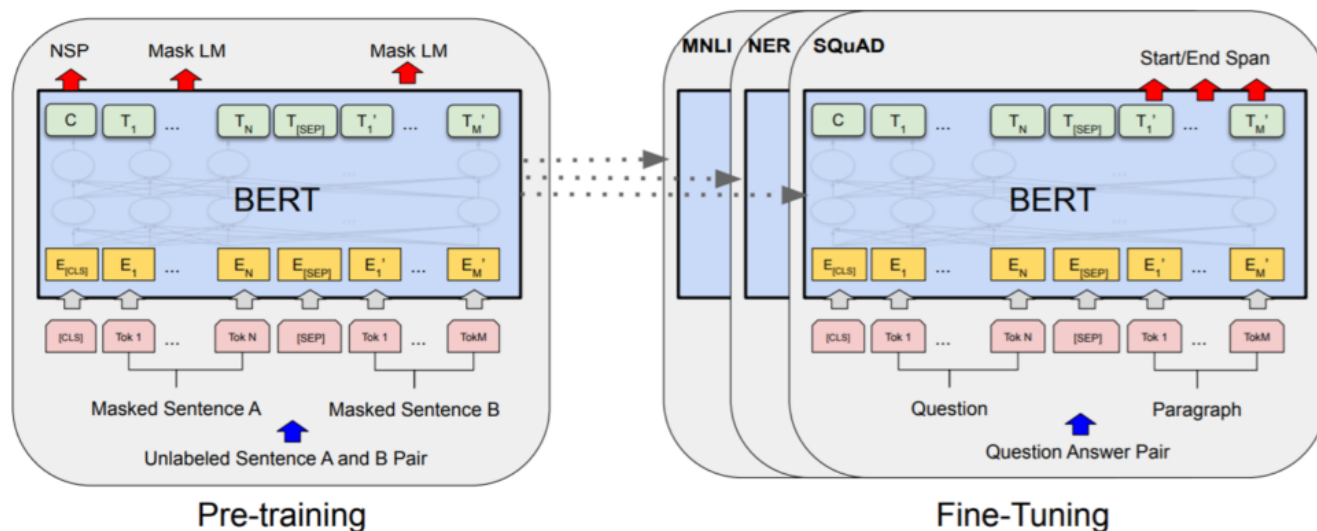
<https://gist.github.com/namakemono/4e2a37375edf7a5e849b0a499897dcbe> to build our LSTM network

We have also optimized our models using the best optimizing strategies to get excellent results.

About BERT

BERT(Bidirectional Encoder Representations from Transformers) is a Transformer-based architecture used for Natural Language processing tasks. Prior to BERT, the state of the art NLP models used LSTM (Long Short Term Memory), which are a certain kind of recurrent neural networks that are slow and not completely bi-directional. This led to the invention of Transformers, that are relatively faster, and can accomodate context-aware word embeddings. BERT stacks a number of encoders to acheive excellent results in question answering, semantic similarity, text classification, etc.

▼ Architecture of BERT



BERT is an encoder stack of transformer architecture. BERT-base has 12 layers in the Encoder stack while BERT-large has 24 layers in the Encoder stack. Bert-base has 110M parameters, while Bert-large has 340M parameters.

BERT trains on two unsupervised tasks simultaneously. These tasks are the Masked Language Model (MLM) and Next Sentence Prediction (NSP). For MLM, It takes in a sentence with random words filled with masks, and the goal is to understand these masked tokens and assign corresponding words by understanding bidirectional context within a sentence.

In the case of NSP, BERT takes in two sentences and returns if the second sentence follows from the first in a manner similar to a binary classification. More specifically, the inputs are fed in the form of pretrained embeddings. The BERT paper combines token embeddings, segment embeddings and position embeddings which are combined and fed as the input. The segment and position embedding are required for temporal understanding, and all these vectors are fed in simultaneously. This helps BERT understand context and thus, using these two BERT understands language in a deeper sense.

BERT Finetuning

Finetuning a BERT model is the task of training the pre-trained model of BERT with the specific problem in hand. This is generally done by training the output layer of our model separately (linear/ LSTM), and adding it to BERT and re-training the entire network.

In this project, we have tried using Linear and LSTM as the output layer.

▼ Common Tasks (Run them before running any model)

▼ 1. Import the required libraries

```
!pip install transformers
```

```
import numpy as np
import pandas as pd
import tensorflow as tf
from transformers import BertForQuestionAnswering
from transformers import BertTokenizer
import transformers
```

Collecting transformers

Downloading <https://files.pythonhosted.org/packages/3a/83/e74092e7f24a08d751aa59b37a9>
|██| 1.3MB 19.5MB/s

Requirement already satisfied: packaging in /usr/local/lib/python3.6/dist-packages (from

Requirement already satisfied: dataclasses; python_version < "3.7" in /usr/local/lib/py

Requirement already satisfied: filelock in /usr/local/lib/python3.6/dist-packages (from

Collecting sacremoses

Downloading <https://files.pythonhosted.org/packages/7d/34/09d19aff26edcc8eb2a01bed8e9>
|██| 890kB 46.0MB/s

Requirement already satisfied: requests in /usr/local/lib/python3.6/dist-packages (from

Collecting tokenizers==0.9.3

Downloading <https://files.pythonhosted.org/packages/4c/34/b39eb9994bc3c999270b69c9eea>
|██| 2.9MB 53.3MB/s

Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.6/dist-packa

Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.6/dist-packages (fr

Requirement already satisfied: protobuf in /usr/local/lib/python3.6/dist-packages (from

Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from tr

Collecting sentencepiece==0.1.91

Downloading <https://files.pythonhosted.org/packages/d4/a4/d0a884c4300004a78cca907a6ff>
|██| 1.1MB 43.7MB/s

Requirement already satisfied: pyparsing>=2.0.2 in /usr/local/lib/python3.6/dist-packag

Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from pack

Requirement already satisfied: click in /usr/local/lib/python3.6/dist-packages (from sa

Requirement already satisfied: joblib in /usr/local/lib/python3.6/dist-packages (from s

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.6/dist-pack

Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/li

Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.6/dist-packages (

Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.6/dist-packa

Requirement already satisfied: setuptools in /usr/local/lib/python3.6/dist-packages (fr

Building wheels for collected packages: sacremoses

Building wheel for sacremoses (setup.py) ... done

Created wheel for sacremoses: filename=sacremoses-0.0.43-cp36-none-any.whl size=89325

Stored in directory: /root/.cache/pip/wheels/29/3c/fd/7ce5c3f0666dab31a50123635e6fb5e

Successfully built sacremoses

Installing collected packages: sacremoses, tokenizers, sentencepiece, transformers

Successfully installed sacremoses-0.0.43 sentencepiece-0.1.91 tokenizers-0.9.3 transfor

▼ 2. Download the SNLI dataset

In this step we download the SNLI dataset extract it into the train validation and test csv files

```
max_length = 128 # Maximum length of input sentence to the model.
```

```
batch_size = 32
```

```
epochs = 2
```

```
# Labels in our dataset.
```

```
labels = ["contradiction", "entailment", "neutral"]
```

```
!curl -LO https://raw.githubusercontent.com/MohamadMerchant/SNLI/master/data.tar.gz
```

```
!tar -xvzf data.tar.gz
```

```

      % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload   Total   Spent    Left   Speed
100 11.1M  100 11.1M    0     0  31.4M      0  --:--:-- --:--:-- --:--:-- 31.3M
SNLI_Corpus/
SNLI_Corpus/snli_1.0_dev.csv
SNLI_Corpus/snli_1.0_train.csv
SNLI_Corpus/snli_1.0_test.csv
```

▼ 3. Load the datasets into pandas dataframes

```
train_df = pd.read_csv("SNLI_Corpus/snli_1.0_train.csv", nrows=100000)
```

```
valid_df = pd.read_csv("SNLI_Corpus/snli_1.0_dev.csv")
```

```
test_df = pd.read_csv("SNLI_Corpus/snli_1.0_test.csv")
```

```
# Shape of the data
```

```
print(f"Total train samples : {train_df.shape[0]}")
```

```
print(f"Total validation samples: {valid_df.shape[0]}")
```

```
print(f"Total test samples: {valid_df.shape[0]}")
```

```
print(f"Sentence1: {train_df.loc[1, 'sentence1']}")
```

```
print(f"Sentence2: {train_df.loc[1, 'sentence2']}")
```

```
print(f"Similarity: {train_df.loc[1, 'similarity']}")
```

```
Total train samples : 100000
```

```
Total validation samples: 10000
```

```
Total test samples: 10000
```

```
Sentence1: A person on a horse jumps over a broken down airplane.
```

```
Sentence2: A person is at a diner, ordering an omelette.
```

```
Similarity: contradiction
```

▼ 4. Clean the data

```
print("Number of missing values")
```

```

print(train_df.isnull().sum())
train_df.dropna(axis=0, inplace=True)
print("Train Target Distribution")
print(train_df.similarity.value_counts())
print("Validation Target Distribution")
print(valid_df.similarity.value_counts())

```

```

Number of missing values
similarity      0
sentence1      0
sentence2       3
dtype: int64
Train Target Distribution
entailment      33384
contradiction    33310
neutral         33193
-               110
Name: similarity, dtype: int64
Validation Target Distribution
entailment      3329
contradiction    3278
neutral         3235
-               158
Name: similarity, dtype: int64

```

```

train_df = (
    train_df[train_df.similarity != "-"]
    .sample(frac=1.0, random_state=42)
    .reset_index(drop=True)
)
valid_df = (
    valid_df[valid_df.similarity != "-"]
    .sample(frac=1.0, random_state=42)
    .reset_index(drop=True)
)

```

▼ 5. Create labels for the train, validation and test datasets

```

train_df["label"] = train_df["similarity"].apply(
    lambda x: 0 if x == "contradiction" else 1 if x == "entailment" else 2
)
y_train = tf.keras.utils.to_categorical(train_df.label, num_classes=3)

valid_df["label"] = valid_df["similarity"].apply(
    lambda x: 0 if x == "contradiction" else 1 if x == "entailment" else 2
)
y_val = tf.keras.utils.to_categorical(valid_df.label, num_classes=3)

test_df["label"] = test_df["similarity"].apply(
    lambda x: 0 if x == "contradiction" else 1 if x == "entailment" else 2
)

```

```
)
y_test = tf.keras.utils.to_categorical(test_df.label, num_classes=3)
```

▼ 6. Sentence similarity verification method

```
def is_similar(sentence1, sentence2, model, tokenizer):
    """
    Takes a sentence1 and checks if sentence2 is symantically similar to sentence1.
    """
    sent = [sentence1,sentence2]

    encoded = tokenizer([sent], return_tensors='pt',add_special_tokens=True,
                        max_length=max_length,
                        return_attention_mask=True,
                        return_token_type_ids=True,
                        pad_to_max_length=True,
                        )

    input_ids = np.array(encoded["input_ids"], dtype="int32")
    attention_masks = np.array(encoded["attention_mask"], dtype="int32")
    token_type_ids = np.array(encoded["token_type_ids"], dtype="int32")

    x_train = [input_ids, attention_masks, token_type_ids]
    # y_train = tf.keras.utils.to_categorical(train_df[0].label, num_classes=3)

    y_pred = np.array(model.predict(x_train))[0]
    print(y_pred)
    idx = np.argmax(y_pred)
    sentiment_labels = ["contradiction", "entailment", "neutral"]
    print(idx)
    print(sentiment_labels[idx])
    print(y_pred[idx])
```

▼ Model 1 : Plain BERT(small)

▼ 1. Create BERT tokenizer

```
tokenizer = transformers.BertTokenizer.from_pretrained(
    "bert-base-uncased", do_lower_case=True
)
```

▼ 2. Encode the train data using BERT tokenizer

```
max_length = 128
```

```
x_train = train_df[["sentence1", "sentence2"]].values.astype("str")
```

```
encoded = tokenizer(x_train[0:100000].tolist(), return_tensors='pt', add_special_tokens=True,
                    max_length=max_length,
                    return_attention_mask=True,
                    return_token_type_ids=True,
                    pad_to_max_length=True,
                    )
```

```
#Embedding and stuff
```

```
input_ids = np.array(encoded["input_ids"], dtype="int32")
```

```
attention_masks = np.array(encoded["attention_mask"], dtype="int32")
```

```
token_type_ids = np.array(encoded["token_type_ids"], dtype="int32")
```

```
x_train = [input_ids, attention_masks, token_type_ids]
```

```
y_train = tf.keras.utils.to_categorical(train_df[0:100000].label, num_classes=3)
```

```
# Encoded token ids from BERT tokenizer.
```

```
input_ids = tf.keras.layers.Input(
    shape=(max_length,), dtype=tf.int32, name="input_ids"
)
```

```
# Attention masks indicates to the model which tokens should be attended to.
```

```
attention_masks = tf.keras.layers.Input(
    shape=(max_length,), dtype=tf.int32, name="attention_masks"
)
```

```
# Token type ids are binary masks identifying different sequences in the model.
```

```
token_type_ids = tf.keras.layers.Input(
    shape=(max_length,), dtype=tf.int32, name="token_type_ids"
)
```

```
#Encode the validation set
```

```
encoded_valid = tokenizer(valid_df[["sentence1", "sentence2"]].values.astype("str").tolist(),
                           add_special_tokens=True,
                           max_length=max_length,
                           return_attention_mask=True,
                           return_token_type_ids=True,
                           pad_to_max_length=True,
                           )
```

```
input_ids_valid = np.array(encoded_valid["input_ids"], dtype="int32")
```

```
attention_masks_valid = np.array(encoded_valid["attention_mask"], dtype="int32")
```

```
token type ids valid = np.array(encoded_valid["token type ids"], dtype="int32")
```



```

x_valid = [input_ids_valid, attention_masks_valid, token_type_ids_valid]
y_valid = tf.keras.utils.to_categorical(valid_df.label, num_classes=3)

valid_data = x_valid,y_valid

```

Truncation was not explicitly activated but `max_length` is provided a specific value,
 /usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022: FutureWarning,

▼ 3. Load the BERT pre-trained model

```

# Loading pretrained BERT model.
bert_model = transformers.TFBertModel.from_pretrained("bert-base-uncased")

# Freeze the BERT model to reuse the pretrained features without modifying them.
bert_model.trainable = False

sequence_output, pooled_output = bert_model(
    input_ids, attention_mask=attention_masks, token_type_ids=token_type_ids
)

sequence_output = tf.keras.layers.Flatten()(sequence_output)
output = tf.keras.layers.Dense(3, activation="softmax")(sequence_output)
model = tf.keras.models.Model(
    inputs=[input_ids, attention_masks, token_type_ids], outputs=output
)

model = tf.keras.models.Model(
    inputs=[input_ids, attention_masks, token_type_ids], outputs=output
)

model.compile(
    optimizer=tf.keras.optimizers.Adam(),
    loss="categorical_crossentropy",
    metrics=["acc"],
)
model.summary()

```

Downloading: 100%

433/433 [00:00<00:00, 3.75kB/s]

Downloading: 100%

536M/536M [00:12<00:00, 43.5MB/s]

Some layers from the model checkpoint at bert-base-uncased were not used when initializing - This IS expected if you are initializing TFBertModel from the checkpoint of a model trained on a different task. - This IS NOT expected if you are initializing TFBertModel from the checkpoint of a model trained on the same task. All the layers of TFBertModel were initialized from the model checkpoint at bert-base-uncased. If your task is similar to the task the model of the checkpoint was trained on, you can initialize the model with the same checkpoint.

Layer (type)	Output Shape	Param #	Connected to
input_ids (InputLayer)	[(None, 128)]	0	
attention_masks (InputLayer)	[(None, 128)]	0	
token_type_ids (InputLayer)	[(None, 128)]	0	
tf_bert_model (TFBertModel)	((None, 128, 768), (None, 128, 768))	109482240	input_ids[0][0] attention_masks[0][0] token_type_ids[0][0]

▼ 4. Train the model

```
training_history = model.fit(x_train,y_train,validation_data=valid_data,epochs=2,batch_size=16)
```

Epoch 1/2

6243/6243 [=====] - 1093s 175ms/step - loss: 4.1711 - acc: 0.5

Epoch 2/2

6243/6243 [=====] - 1104s 177ms/step - loss: 4.3283 - acc: 0.5

▼ 5. Save the model (Optional)

```
model.save_weights('BERT_Plain.h5')
```

▼ 6. Verify the model performance

```
is_similar("Charlie went to the cycle shop on Sunday","Charlie was resting at home on Sunday")
```

```
/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022: FutureWarning,
[5.3576505e-01 2.0653822e-11 4.6423498e-01]
0
```

```
contradiction
0.53576505
```

▼ 7. Test the model

```
encoded = tokenizer(test_df[["sentence1", "sentence2"]].values.astype("str").tolist(), return_
    add_special_tokens=True,
    max_length=max_length,
    return_attention_mask=True,
    return_token_type_ids=True,
    pad_to_max_length=True,
    )
```

```
input_ids = np.array(encoded["input_ids"], dtype="int32")
attention_masks = np.array(encoded["attention_mask"], dtype="int32")
token_type_ids = np.array(encoded["token_type_ids"], dtype="int32")
```

```
x_test = [input_ids, attention_masks, token_type_ids]
y_test = tf.keras.utils.to_categorical(test_df.label, num_classes=3)
```

```
model.evaluate(x_test,y_test, verbose=1)
```

```
/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022: FutureWarning,
313/313 [=====] - 89s 283ms/step - loss: 4.4134 - acc: 0.6278
[4.413364887237549, 0.6277999877929688]
```

▼ 8. Metrics

```
from sklearn import metrics
import matplotlib.pyplot as plt
```

```
y_pred = model.predict(x_test, verbose=1)
```


```
print("Confusion Matrix:\n")
print(metrics.confusion_matrix(np.argmax(y_test, axis=1), np.argmax(y_pred, axis=1)))
```

```
loss_train = training_history.history['loss']
loss_val = training_history.history['val_loss']
epochs = np.arange(1,3)
plt.plot(epochs, np.array(loss_train), 'g', label='Training loss')
plt.plot(epochs, np.array(loss_val), 'b', label='validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
```

```
plt.ylabel('Loss')
plt.legend()
plt.show()
```

313/313 [=====] - 89s 284ms/step
Confusion Matrix:

```
[[2414  612  211]
 [ 475 2704  189]
 [1244  991 1160]]
```



▼ Model 2 : LSTM

▼ 1. Create word embeddings

```
!wget -c "https://s3.amazonaws.com/dl4j-distribution/GoogleNews-vectors-negative300.bin.gz"

--2020-11-28 10:55:47-- https://s3.amazonaws.com/dl4j-distribution/GoogleNews-vectors-negative300.bin.gz
Resolving s3.amazonaws.com (s3.amazonaws.com)... 52.216.139.53
Connecting to s3.amazonaws.com (s3.amazonaws.com)|52.216.139.53|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1647046227 (1.5G) [application/x-gzip]
Saving to: 'GoogleNews-vectors-negative300.bin.gz'

GoogleNews-vectors- 100%[=====>] 1.53G 81.0MB/s in 22s

2020-11-28 10:56:10 (70.0 MB/s) - 'GoogleNews-vectors-negative300.bin.gz' saved [164704
```

▼ 2. Create the model and run it

```

import numpy as np
import pandas as pd
from gensim.models import KeyedVectors
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.layers import Dense, Input, LSTM, Embedding, Dropout, Activation
from keras.layers.merge import concatenate
from keras.models import Model
from keras.layers.normalization import BatchNormalization
from keras.callbacks import EarlyStopping
from keras.utils import to_categorical

EMBEDDING_FILE = 'GoogleNews-vectors-negative300.bin.gz'
TRAIN_DATA_FILE = "SNLI_Corpus/snli_1.0_train.csv"
TEST_DATA_FILE = "SNLI_Corpus/snli_1.0_test.csv"
MAX_SEQUENCE_LENGTH = 30
MAX_NB_WORDS = 200000
EMBEDDING_DIM = 300
VALIDATION_SPLIT = 0.1
rate_drop_lstm = 0.15 + np.random.rand() * 0.25
rate_drop_dense = 0.15 + np.random.rand() * 0.25

def text_to_tokens(text):
    return text.lower()

def get_label_index_mapping():
    return {"neutral": 0, "contradiction": 1, "entailment": 2, "-": 3}

def create_embedding_matrix(word_index):
    nb_words = min(MAX_NB_WORDS, len(word_index))+1
    word2vec = KeyedVectors.load_word2vec_format(EMBEDDING_FILE, binary=True)
    embedding_matrix = np.zeros((nb_words, EMBEDDING_DIM))
    for word, i in word_index.items():
        if word in word2vec.vocab:
            embedding_matrix[i] = word2vec.word_vec(word)
    return embedding_matrix

def load_data():
    train_df = pd.read_csv(TRAIN_DATA_FILE, sep=",", usecols=["sentence1", "sentence2", "simi"])
    train_df.fillna("", inplace=True)
    sentence1 = train_df["sentence1"].apply(text_to_tokens)
    sentence2 = train_df["sentence2"].apply(text_to_tokens)
    y = train_df["similarity"].map(get_label_index_mapping())
    tokenizer = Tokenizer(num_words=MAX_NB_WORDS)
    tokenizer.fit_on_texts(sentence1 + sentence2)
    sequences1 = tokenizer.texts_to_sequences(sentence1)
    sequences2 = tokenizer.texts_to_sequences(sentence2)
    X1 = pad_sequences(sequences1, maxlen=MAX_SEQUENCE_LENGTH)

```

```

...
X2 = pad_sequences(sequences2, maxlen=MAX_SEQUENCE_LENGTH)
perm = np.random.permutation(len(X1))
num_train = int(len(X1)*(1-VALIDATION_SPLIT))
train_index = perm[:num_train]
valid_index = perm[num_train:]
X1_train = X1[train_index]
X2_train = X2[train_index]
y_train = y[train_index]
X1_valid = X1[valid_index]
X2_valid = X2[valid_index]
y_valid = y[valid_index]
return (X1_train, X2_train, y_train), (X1_valid, X2_valid, y_valid), tokenizer

def StaticEmbedding(embedding_matrix):
    input_dim, output_dim = embedding_matrix.shape
    return Embedding(input_dim,
                      output_dim,
                      weights=[embedding_matrix],
                      input_length=MAX_SEQUENCE_LENGTH,
                      trainable=False)

def entail(feat1, feat2, num_dense=300):
    x = concatenate([feat1, feat2])
    x = Dropout(rate_drop_dense)(x)
    x = BatchNormalization()(x)
    x = Dense(num_dense, activation="relu")(x)
    x = Dropout(rate_drop_dense)(x)
    x = BatchNormalization()(x)
    return x

def build_model(output_dim, embedding_matrix, num_lstm=300):
    sequence1_input = Input(shape=(MAX_SEQUENCE_LENGTH,), dtype='int32')
    sequence2_input = Input(shape=(MAX_SEQUENCE_LENGTH,), dtype='int32')

    # Embedding
    embed = StaticEmbedding(embedding_matrix)
    embedded_sequences1 = embed(sequence1_input)
    embedded_sequences2 = embed(sequence2_input)

    # Encoding
    encode = LSTM(num_lstm, dropout=rate_drop_lstm, recurrent_dropout=rate_drop_lstm)
    feat1 = encode(embedded_sequences1)
    feat2 = encode(embedded_sequences2)

    x = entail(feat1, feat2)
    preds = Dense(output_dim, activation='softmax')(x)
    model = Model(inputs=[sequence1_input, sequence2_input], outputs=preds)
    return model

def run():
    num_class = len(get_label_index_mapping())

```

```
(X1_train, X2_train, y_train), (X1_valid, X2_valid, y_valid), tokenizer = load_data()
Y_train, Y_valid = to_categorical(y_train, num_class), to_categorical(y_valid, num_class)
embedding_matrix = create_embedding_matrix(tokenizer.word_index)
model = build_model(output_dim=num_class, embedding_matrix=embedding_matrix)
model.compile(loss='categorical_crossentropy', optimizer='nadam', metrics=['acc'])
early_stopping = EarlyStopping(monitor='val_loss', patience=10)
hist = model.fit([X1_train, X2_train], Y_train,
                 validation_data=([X1_valid, X2_valid], Y_valid),
                 epochs=20, batch_size=2048, shuffle=True,
                 callbacks=[early_stopping])
```

```
if __name__ == "__main__":
    run()
```

```
WARNING:tensorflow:Layer lstm will not use cuDNN kernel since it doesn't meet the cuDNN
Epoch 1/20
242/242 [=====] - 88s 363ms/step - loss: 1.0785 - acc: 0.5303
Epoch 2/20
242/242 [=====] - 88s 365ms/step - loss: 0.8565 - acc: 0.6141
Epoch 3/20
242/242 [=====] - 88s 365ms/step - loss: 0.8016 - acc: 0.6476
Epoch 4/20
242/242 [=====] - 89s 367ms/step - loss: 0.7619 - acc: 0.6711
Epoch 5/20
242/242 [=====] - 89s 369ms/step - loss: 0.7321 - acc: 0.6873
Epoch 6/20
242/242 [=====] - 89s 368ms/step - loss: 0.7112 - acc: 0.6989
Epoch 7/20
242/242 [=====] - 89s 367ms/step - loss: 0.6943 - acc: 0.7080
Epoch 8/20
242/242 [=====] - 89s 368ms/step - loss: 0.6797 - acc: 0.7157
Epoch 9/20
242/242 [=====] - 89s 369ms/step - loss: 0.6686 - acc: 0.7213
Epoch 10/20
242/242 [=====] - 89s 368ms/step - loss: 0.6586 - acc: 0.7265
Epoch 11/20
242/242 [=====] - 89s 367ms/step - loss: 0.6501 - acc: 0.7312
Epoch 12/20
242/242 [=====] - 89s 369ms/step - loss: 0.6425 - acc: 0.7343
Epoch 13/20
242/242 [=====] - 89s 369ms/step - loss: 0.6362 - acc: 0.7376
Epoch 14/20
242/242 [=====] - 89s 369ms/step - loss: 0.6285 - acc: 0.7421
Epoch 15/20
242/242 [=====] - 89s 368ms/step - loss: 0.6228 - acc: 0.7439
Epoch 16/20
242/242 [=====] - 89s 367ms/step - loss: 0.6191 - acc: 0.7460
Epoch 17/20
242/242 [=====] - 89s 367ms/step - loss: 0.6140 - acc: 0.7488
Epoch 18/20
242/242 [=====] - 89s 368ms/step - loss: 0.6080 - acc: 0.7508
Epoch 19/20
242/242 [=====] - 89s 369ms/step - loss: 0.6054 - acc: 0.7535
Epoch 20/20
242/242 [=====] - 90s 370ms/step - loss: 0.6009 - acc: 0.7548
```

▼ Model 3 : BERT-large pre-trained on SQuAD

▼ 1. Create BERT-large tokenizer

```
tokenizer_b1 = BertTokenizer.from_pretrained('bert-large-uncased-whole-word-masking-finetuned-squad-uncased')
```

Downloading: 100%

232k/232k [00:00<00:00, 2.02MB/s]

▼ 2. Encode the train data using BERT-large tokenizer

```
max_length = 128
```

```
x_train = train_df[["sentence1", "sentence2"]].values.astype("str")
```

```
encoded = tokenizer_b1(x_train[0:100000].tolist(), return_tensors='pt', add_special_tokens=True,
                        max_length=max_length,
                        return_attention_mask=True,
                        return_token_type_ids=True,
                        pad_to_max_length=True,
                        )
```

```
#Embedding and stuff
```

```
input_ids = np.array(encoded["input_ids"], dtype="int32")
```

```
attention_masks = np.array(encoded["attention_mask"], dtype="int32")
```

```
token_type_ids = np.array(encoded["token_type_ids"], dtype="int32")
```

```
x_train = [input_ids, attention_masks, token_type_ids]
```

```
y_train = tf.keras.utils.to_categorical(train_df[0:100000].label, num_classes=3)
```

```
# Encoded token ids from BERT tokenizer.
```

```
input_ids = tf.keras.layers.Input(
    shape=(max_length,), dtype=tf.int32, name="input_ids"
)
```

```
# Attention masks indicates to the model which tokens should be attended to.
```

```
attention_masks = tf.keras.layers.Input(
    shape=(max_length,), dtype=tf.int32, name="attention_masks"
)
```

```
# Token type ids are binary masks identifying different sequences in the model.
```

```
token_type_ids = tf.keras.layers.Input(
    shape=(max_length,), dtype=tf.int32, name="token type ids"
)
```



```

)

#Encode the validation set
encoded_valid = tokenizer_b1(valid_df[["sentence1", "sentence2"]].values.astype("str").tolist
    add_special_tokens=True,
    max_length=max_length,
    return_attention_mask=True,
    return_token_type_ids=True,
    pad_to_max_length=True,
)

input_ids_valid = np.array(encoded_valid["input_ids"], dtype="int32")
attention_masks_valid = np.array(encoded_valid["attention_mask"], dtype="int32")
token_type_ids_valid = np.array(encoded_valid["token_type_ids"], dtype="int32")

x_valid = [input_ids_valid, attention_masks_valid, token_type_ids_valid]
y_valid = tf.keras.utils.to_categorical(valid_df.label, num_classes=3)

valid_data = x_valid,y_valid

Truncation was not explicitly activated but `max_length` is provided a specific value,
/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022: FutureWarning,

```

▼ 3. Load the BERT-large pre-trained model

```

# Loading pretrained BERT model.
bert_model = transformers.TFBertModel.from_pretrained("bert-large-uncased-whole-word-masking-

# Freeze the BERT model to reuse the pretrained features without modifying them.
bert_model.trainable = False

sequence_output, pooled_output = bert_model(
    input_ids, attention_mask=attention_masks, token_type_ids=token_type_ids
)

# Add trainable layers on top of frozen layers to adapt the pretrained features on the new data
bi_lstm = tf.keras.layers.Bidirectional(
    tf.keras.layers.LSTM(64, return_sequences=True)
)(sequence_output)

# Applying hybrid pooling approach to bi_lstm sequence output.
avg_pool = tf.keras.layers.GlobalAveragePooling1D()(bi_lstm)
max_pool = tf.keras.layers.GlobalMaxPooling1D()(bi_lstm)
concat = tf.keras.layers.concatenate([avg_pool, max_pool])
dropout = tf.keras.layers.Dropout(0.3)(concat)

```

```
output = tf.keras.layers.Dense(3, activation="softmax")(dropout)

model_b1 = tf.keras.models.Model(
    inputs=[input_ids, attention_masks, token_type_ids], outputs=output
)

model_b1.compile(
    optimizer=tf.keras.optimizers.Adam(),
    loss="categorical_crossentropy",
    metrics=["acc"],
)
model_b1.summary()
```

▼ 4. Train the model

```
training_history = model_b1.fit(x_train,y_train,validation_data=valid_data,epochs=2,batch_size=128)
```

```
Epoch 1/2
6243/6243 [=====] - 3323s 532ms/step - loss: 0.7092 - acc: 0.6
Epoch 2/2
6243/6243 [=====] - 3329s 533ms/step - loss: 0.6277 - acc: 0.7
```

Layer (type)	Output Shape	Param #	Connected to
--------------	--------------	---------	--------------

▼ 5. Finetune the model

```
# Unfreeze the bert_model.
bert_model.trainable = True
# Recompile the model to make the change effective.
model_b1.compile(
    optimizer=tf.keras.optimizers.Adam(1e-5),
    loss="categorical_crossentropy",
    metrics=["accuracy"],
)
model_b1.summary()
```

```
finetuning_history = model_b1.fit(x_train,y_train, validation_data=valid_data, epochs=2,batch_size=128)
```

Model: "functional_1"

Layer (type)	Output Shape	Param #	Connected to
input_ids (InputLayer)	[(None, 128)]	0	
attention_masks (InputLayer)	[(None, 128)]	0	
token_type_ids (InputLayer)	[(None, 128)]	0	
tf_bert_model (TFBertModel)	((None, 128, 1024),	335141888	input_ids[0][0] attention_masks[0][0] token_type_ids[0][0]
bidirectional (Bidirectional)	(None, 128, 128)	557568	tf_bert_model[0][0]
global_average_pooling1d (GlobalAveragePooling1D)	(None, 128)	0	bidirectional[0][0]
global_max_pooling1d (GlobalMaxPooling1D)	(None, 128)	0	bidirectional[0][0]
concatenate (Concatenate)	(None, 256)	0	global_average_pooling1d[0] global_max_pooling1d[0]
dropout_73 (Dropout)	(None, 256)	0	concatenate[0][0]

```

dense (Dense)                (None, 3)                771                dropout_73[0][0]
=====
Total params: 335,700,227
Trainable params: 335,700,227
Non-trainable params: 0
=====
Epoch 1/2
WARNING:tensorflow:Gradients do not exist for variables ['tf_bert_model/bert/pooler/den
WARNING:tensorflow:Gradients do not exist for variables ['tf_bert_model/bert/pooler/den
WARNING:tensorflow:Gradients do not exist for variables ['tf_bert_model/bert/pooler/den
WARNING:tensorflow:Gradients do not exist for variables ['tf_bert_model/bert/pooler/den
6243/6243 [=====] - 9002s 1s/step - loss: 0.4614 - accuracy: 0
Epoch 2/2
6243/6243 [=====] - 8998s 1s/step - loss: 0.3380 - accuracy: 0

```

► 6. Save the model (Optional)

```
[ ] ↳ 1 cell hidden
```

▼ 7. Verify the model performance

y to sample the region's exquisite wines.", "Be sure to make time for a Tuscan wine-tasting ex

```

[0.01231802 0.7330076 0.2546743 ]
1
entailment
0.7330076
/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022: Fu
FutureWarning,

```

▼ 8. Test the model

```

encoded = tokenizer_b1(test_df[["sentence1", "sentence2"]].values.astype("str").tolist(), ret
add_special_tokens=True,
max_length=max_length,
return_attention_mask=True,
return_token_type_ids=True,
pad_to_max_length=True,
)

```

```

input_ids = np.array(encoded["input_ids"], dtype="int32")
attention_masks = np.array(encoded["attention_mask"], dtype="int32")
token_type_ids = np.array(encoded["token_type_ids"], dtype="int32")

```

```

x_test = [input_ids, attention_masks, token_type_ids]
y_test = tf.keras.utils.to_categorical(test_df.label, num_classes=3)

model_b1.evaluate(x_test,y_test, verbose=1)

/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022: FutureWarning,
313/313 [=====] - 281s 898ms/step - loss: 0.3264 - accuracy: 0
[0.32644084095954895, 0.8827000260353088]

```

▼ 9. Metrics

```

from sklearn import metrics
import matplotlib.pyplot as plt

y_pred = model_b1.predict(x_test, verbose=1)

print("Confusion Matrix:\n")
print(metrics.confusion_matrix(np.argmax(y_test, axis=1), np.argmax(y_pred, axis=1)))

```

```

loss_train = training_history.history['loss']
loss_val = training_history.history['val_loss']
epochs = np.arange(1,3)
plt.plot(epochs, np.array(loss_train), 'g', label='Training loss')
plt.plot(epochs, np.array(loss_val), 'b', label='validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()

```

```



loss_train = finetuning_history.history['loss']
loss_val = finetuning_history.history['val_loss']
epochs = np.arange(1,3)
plt.plot(epochs, np.array(loss_train), 'g', label='Training loss')
plt.plot(epochs, np.array(loss_val), 'b', label='validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()

```

313/313 [=====] - 283s 903ms/step

Confusion Matrix:

```
[[2938  91  208]
 [  62 3058  248]
 [ 224  340 2831]]
```



▼ Model 4 : Plain BERT(small) with Bi-Directional LSTM for finetuning

▼ 1. Create BERT tokenizer

```
tokenizer = transformers.BertTokenizer.from_pretrained(
    "bert-base-uncased", do_lower_case=True
)
```

Downloading: 100%

232k/232k [00:00<00:00, 2.13MB/s]

▼ 2. Encode the train data using BERT tokenizer

```

max_length = 128

x_train = train_df[["sentence1", "sentence2"]].values.astype("str")

encoded = tokenizer(x_train[0:100000].tolist(), return_tensors='pt', add_special_tokens=True,
                    max_length=max_length,
                    return_attention_mask=True,
                    return_token_type_ids=True,
                    pad_to_max_length=True,
                    )

#Embedding and stuff
input_ids = np.array(encoded["input_ids"], dtype="int32")
attention_masks = np.array(encoded["attention_mask"], dtype="int32")
token_type_ids = np.array(encoded["token_type_ids"], dtype="int32")

x_train = [input_ids, attention_masks, token_type_ids]
y_train = tf.keras.utils.to_categorical(train_df[0:100000].label, num_classes=3)

# Encoded token ids from BERT tokenizer.
input_ids = tf.keras.layers.Input(
    shape=(max_length,), dtype=tf.int32, name="input_ids"
)
# Attention masks indicates to the model which tokens should be attended to.
attention_masks = tf.keras.layers.Input(
    shape=(max_length,), dtype=tf.int32, name="attention_masks"
)
# Token type ids are binary masks identifying different sequences in the model.
token_type_ids = tf.keras.layers.Input(
    shape=(max_length,), dtype=tf.int32, name="token_type_ids"
)

#Encode the validation set
encoded_valid = tokenizer(valid_df[["sentence1", "sentence2"]].values.astype("str").tolist(),
                           add_special_tokens=True,
                           max_length=max_length,
                           return_attention_mask=True,
                           return_token_type_ids=True,
                           pad_to_max_length=True,
                           )


input_ids_valid = np.array(encoded_valid["input_ids"], dtype="int32")
attention_masks_valid = np.array(encoded_valid["attention_mask"], dtype="int32")
token_type_ids_valid = np.array(encoded_valid["token_type_ids"], dtype="int32")

```

```
x_valid = [input_ids_valid, attention_masks_valid, token_type_ids_valid]
y_valid = tf.keras.utils.to_categorical(valid_df.label, num_classes=3)
```

```
valid_data = x_valid,y_valid
```

```
/usr/local/lib/python3.6/dist-packages/transformers/tokenization_utils_base.py:2022: FutureWarning,
```



▼ 3. Load the BERT pre-trained model

```
# Loading pretrained BERT model.
```

```
bert_model = transformers.TFBertModel.from_pretrained("bert-base-uncased")
```

```
# Freeze the BERT model to reuse the pretrained features without modifying them.
```

```
bert_model.trainable = False
```

```
sequence_output, pooled_output = bert_model(
    input_ids, attention_mask=attention_masks, token_type_ids=token_type_ids
)
```

```
# Add trainable layers on top of frozen layers to adapt the pretrained features on the new data
```

```
bi_lstm = tf.keras.layers.Bidirectional(
    tf.keras.layers.LSTM(256, return_sequences=True)
)(sequence_output)
```

```
# Applying hybrid pooling approach to bi_lstm sequence output.
```

```
avg_pool = tf.keras.layers.GlobalAveragePooling1D()(bi_lstm)
max_pool = tf.keras.layers.GlobalMaxPooling1D()(bi_lstm)
concat = tf.keras.layers.concatenate([avg_pool, max_pool])
dropout = tf.keras.layers.Dropout(0.2)(concat)
```

```
output = tf.keras.layers.Dense(3, activation="softmax")(dropout)
```

```
model_ft = tf.keras.models.Model(
    inputs=[input_ids, attention_masks, token_type_ids], outputs=output
)
```

```
model_ft.compile(
    optimizer=tf.keras.optimizers.Adam(),
    loss="categorical_crossentropy",
    metrics=["acc"],
)
model_ft.summary()
```


Downloading: 100%

433/433 [00:00<00:00, 13.1kB/s]

Downloading: 100%

536M/536M [00:22<00:00, 23.3MB/s]

Some layers from the model checkpoint at bert-base-uncased were not used when initializing - This IS expected if you are initializing TFBertModel from the checkpoint of a model trained on a different task. This IS NOT expected if you are initializing TFBertModel from the checkpoint of a model trained on the same task. All the layers of TFBertModel were initialized from the model checkpoint at bert-base-uncased. If your task is similar to the task the model of the checkpoint was trained on, you can initialize the model with the same weights. Model: "functional_3"

Layer (type)	Output Shape	Param #	Connected to
input_ids (InputLayer)	[(None, 128)]	0	
attention_masks (InputLayer)	[(None, 128)]	0	
token_type_ids (InputLayer)	[(None, 128)]	0	
tf_bert_model_1 (TFBertModel)	((None, 128, 768), (None, 128, 768))	109482240	input_ids[0][0] attention_masks[0][0] token_type_ids[0][0]
bidirectional_1 (Bidirectional)	(None, 128, 512)	2099200	tf_bert_model_1[0][0]
global_average_pooling1d_1 (GlobalAveragePooling1D)	(None, 512)	0	bidirectional_1[0][0]
global_max_pooling1d_1 (GlobalMaxPooling1D)	(None, 512)	0	bidirectional_1[0][0]
concatenate_1 (Concatenate)	(None, 1024)	0	global_average_pooling1d_1 global_max_pooling1d_1
dropout_111 (Dropout)	(None, 1024)	0	concatenate_1[0][0]
dense_1 (Dense)	(None, 3)	3075	dropout_111[0][0]
Total params: 111,584,515			
Trainable params: 2,102,275			

▼ 4. Train the model

```
training_history = model_ft.fit(x_train,y_train,validation_data=valid_data,epochs=2,batch_size=16)
```

Epoch 1/2

6243/6243 [=====] - 1228s 197ms/step - loss: 0.6517 - acc: 0.7

Epoch 2/2

6243/6243 [=====] - 1222s 196ms/step - loss: 0.5699 - acc: 0.7

▼ 5. Finetune the model

```
# Unfreeze the bert_model.
bert_model.trainable = True
# Recompile the model to make the change effective.
model_ft.compile(
    optimizer=tf.keras.optimizers.Adam(1e-5),
    loss="categorical_crossentropy",
    metrics=["accuracy"],
)
model_ft.summary()

finetuning_history = model_ft.fit(x_train,y_train, validation_data=valid_data, epochs=2,batch
```

Model: "functional_3"

Layer (type)	Output Shape	Param #	Connected to
input_ids (InputLayer)	[(None, 128)]	0	
attention_masks (InputLayer)	[(None, 128)]	0	
token_type_ids (InputLayer)	[(None, 128)]	0	
tf_bert_model_1 (TFBertModel)	((None, 128, 768), (109482240	input_ids[0][0] attention_masks[0][0] token_type_ids[0][0]
bidirectional_1 (Bidirectional)	(None, 128, 512)	2099200	tf_bert_model_1[0][0]
global_average_pooling1d_1 (Glo	(None, 512)	0	bidirectional_1[0][0]
global_max_pooling1d_1 (GlobalM	(None, 512)	0	bidirectional_1[0][0]
concatenate_1 (Concatenate)	(None, 1024)	0	global_average_pooling global_max_pooling1d_1
dropout_111 (Dropout)	(None, 1024)	0	concatenate_1[0][0]
dense_1 (Dense)	(None, 3)	3075	dropout_111[0][0]

Total params: 111,584,515

Trainable params: 111,584,515

Non-trainable params: 0

Epoch 1/2

WARNING:tensorflow:Gradients do not exist for variables ['tf_bert_model_1/bert/pooler/d

WARNING:tensorflow:Gradients do not exist for variables ['tf_bert_model_1/bert/pooler/d

WARNING:tensorflow:Gradients do not exist for variables ['tf_bert_model_1/bert/pooler/d

WARNING:tensorflow:Gradients do not exist for variables ['tf_bert_model_1/bert/pooler/d

6243/6243 [=====] - 3001s 481ms/step - loss: 0.4487 - accuracy

Epoch 2/2

4665/6243 [=====>.....] - ETA: 12:13 - loss: 0.3209 - accuracy: 0.88

▼ 6. Save the model (Optional)

```
model_ft.save_weights('BERT_Finetuned_Experiment1.h5')
```

▼ 7. Verify the model performance

```
is_similar("Charlie went to the cycle shop on Sunday","Charlie was resting at home on Sunday")
```

▼ 8. Test the model

```
encoded = tokenizer(test_df[["sentence1", "sentence2"]].values.astype("str").tolist(), return_
    add_special_tokens=True,
    max_length=max_length,
    return_attention_mask=True,
    return_token_type_ids=True,
    pad_to_max_length=True,
    )
```

```
input_ids = np.array(encoded["input_ids"], dtype="int32")
attention_masks = np.array(encoded["attention_mask"], dtype="int32")
token_type_ids = np.array(encoded["token_type_ids"], dtype="int32")
```

```
x_test = [input_ids, attention_masks, token_type_ids]
y_test = tf.keras.utils.to_categorical(test_df.label, num_classes=3)
```

```
model_ft.evaluate(x_test,y_test, verbose=1)
```

▼ 9. Metrics

```
from sklearn import metrics
import matplotlib.pyplot as plt
```

```
y_pred = model_ft.predict(x_test, verbose=1)
```

```
print("Confusion Matrix:\n")
print(metrics.confusion_matrix(np.argmax(y_test, axis=1), np.argmax(y_pred, axis=1)))
```

```
loss_train = training_history.history['loss']
loss_val = training_history.history['val_loss']
epochs = np.arange(1,3)
plt.plot(epochs, np.array(loss_train), 'g', label='Training loss')
```

```
plt.plot(epochs, np.array(loss_val), 'b', label='validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

```
loss_train = finetuning_history.history['loss']
loss_val = finetuning_history.history['val_loss']
epochs = np.arange(1,2)
plt.plot(epochs, np.array(loss_train), 'g', label='Training loss')
plt.plot(epochs, np.array(loss_val), 'b', label='validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

▼ Experiments

Experiment 1 : Plain Bert(Small) without finetuning

- Dense output layer
- bert-base-uncased tokenizer
- 100000 training examples

Results:

After Training:

```
loss: 4.1711 - acc: 0.5355
val_loss: 3.8692 - val_acc: 0.6167
```

Test data:

```
loss: 4.4134 - acc: 0.6278
```

We can see that this model performs poorly.

Experiment 2 : LSTM

- 20 epochs
- word2vec for embeddings
- 300 layers in the dense network

Results:

After Training:

```
loss: 0.6009 - acc: 0.7548  
val_loss: 0.5742 - val_acc: 0.7664
```

We can see that the training and validation sets give bad results.

Experiment 3 : BERT-large pre-trained on SQuAD

- pre-trained on SQuAD
- 100000 training examples
- 64 node Bi-directional LSTM in output layer
- Dropout = 0.3

Results:

After Training:

```
Epoch 1:  
loss: 0.7092 - acc: 0.6970  
val_loss: 0.5043 - val_acc: 0.8027
```

```
Epoch 2:  
loss: 0.6277 - acc: 0.7438  
val_loss: 0.4929 - val_acc: 0.8092
```

After Finetuning:

```
Epoch 1:  
loss: 0.4614 - accuracy: 0.8273  
val_loss: 0.3286 - val_accuracy: 0.8795
```

```
Epoch 2:  
loss: 0.3380 - accuracy: 0.8799  
val_loss: 0.2995 - val_accuracy: 0.8916
```

Test data:

```
loss: 0.3264 - accuracy: 0.8827
```

Experiment 4 : Plain Bert(Small) with Bi-directional LSTM for finetuning

- 128 nodes in output LSTM layer
- dropout : 0.35

Results:

After Training:

```
loss: 0.5971 - acc: 0.7559  
val_loss: 0.5043 - val_acc: 0.8027
```

After Finetuning:

```
loss: 0.4692 - accuracy: 0.8173  
val_loss: 0.3623 - val_accuracy: 0.8683
```

Test data:

```
loss: 0.3819 - accuracy: 0.8587
```

Experiment 5 : Plain Bert(Small) with Bi-directional LSTM for finetuning

- 256 nodes in output LSTM layer
- Decreased dropout rate to 0.2

Results:

After Training:

```
Epoch 1:  
loss: 0.6517 - acc: 0.7248  
val_loss: 0.5049 - val_acc: 0.8002
```

```
Epoch 2:  
loss: 0.5699 - acc: 0.7681  
val_loss: 0.5088 - val_acc: 0.7995
```

After Finetuning:

```
Epoch 1:  
loss: 0.4487 - accuracy: 0.8261
```

```
val_loss: 0.3707 - val_accuracy: 0.8625
```

Epoch 2:

```
loss: 0.3209 - accuracy: 0.8822
```

▼ Conclusion

From the above experiments, we observed that the Experiment 3 : BERT-large pre-trained on SQuAD and finetuned with Bi-directional LSTM was performing the best, closely followed by Experiment 4 and 5 on Plain BERT(small) with Bidirectional LSTM for finetuning. From the training and validation set plots for these models, we could observe that the models are underfitting.

We have used the Model 4 in our two games (SemanticSimilarityGame.ipynb and ReadingComprehensionGame.ipynb) , and achieved a great gameplay experience.

References

- [1] BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding Jacob Devlin, Ming-Wei Chang, Kenton Lee, Kristina Toutanova <https://arxiv.org/pdf/1810.04805.pdf>
- [2] Semantic Similarity with BERT Author: Mohamad Merchant https://keras.io/examples/nlp/semantic_similarity_with_bert/
- [3] LSTM with SNLI dataset : <https://gist.github.com/namakemono/4e2a37375edf7a5e849b0a499897dcbe>
- [4] Stanford SNLI Corpus : <https://nlp.stanford.edu/projects/snli/>
- [5] Saving and Loading Keras models: https://www.tensorflow.org/guide/keras/save_and_serialize
- [6] BERT Explained: State of the art language model for NLP - Rani Horev <https://towardsdatascience.com/bert-explained-state-of-the-art-language-model-for-nlp-f8b21a9b6270>
- [7] BERT Neural Network Explained! - <https://www.youtube.com/watch?v=xI0HHN5XKDo>

